

SplitPix — System Design Reference

Version: 4.0 · **Date:** 2026-08-16 · **Status:** current

Scope: the design of the SplitPix group-expense service as implemented — components, boundaries, invariants, and the reasoning behind each load-bearing decision.

Table of contents

1. Purpose and scope
 2. Design goals and constraints
 3. System overview
 4. Architectural decisions
 5. File and directory organization
 6. Core types and classes
 7. Interfaces and boundaries
 8. Invariants and assumptions
 9. Known gaps and future work
 10. The settlement optimization engine
 11. Balance provenance
 12. Against a conventional expense-sharing design
- Appendix A: glossary
-

1. Purpose and scope

1.1 The problem

A group shares costs: one person pays for dinner, another buys groceries, a third covers the taxi. Settling up by hand means reconstructing who paid what, computing each person's net position, and agreeing on who transfers money to whom. In Brazil the transfer itself is easy — Pix is instant and free — so the friction is entirely in the accounting and the coordination, not the payment.

SplitPix records the shared expenses, derives each participant's net balance, and turns the resulting web of debts into a settlement plan — for each debtor, a recipient, an exact amount, and the recipient's Pix key. The word *a* in "a settlement plan" is doing real work: many different transfer graphs leave every participant at exactly zero, and they differ in ways a group can care about — how many payments there are, and whether money moves between people who never transacted. SplitPix treats the choice among those graphs as an explicit optimization problem with selectable strategies (§10), and can explain both any balance (§11) and any plan it proposes.

1.2 What operates it

A single Spring Boot application process plus one PostgreSQL database. There is no scheduler, no queue, no background worker, and no second service. A group is created by one person, who shares an invite link; everyone holding that link operates on the same group through either the REST API or the browser UI.

1.3 What it produces

- Net balances per participant, derived on demand (§4.3), each explainable down to its ledger entries (§11).
- Settlement plans under three strategies, generated on demand and never stored (§4.6, §10), each stamped with the ledger revision it was derived from.
- An append-only record of expenses and completed settlements, with dense sequence numbers.

1.4 Explicitly out of scope

SplitPix **does not move money**. It initiates no Pix transfer, holds no bank credentials, integrates with no financial institution, and cannot verify that a payment occurred — marking a settlement complete is a human assertion. Also out of scope: user accounts and authentication (§4.7), multi-currency, interest or fees, expense editing or deletion (§4.8), participant removal, and public group discovery.

2. Design goals and constraints

Ranked. Each subsequent decision in §4 traces to a goal here; where two goals conflict, the higher-ranked one wins.

G1. Accounting correctness is non-negotiable. Money must never be created or destroyed by the system. The sum of a group's balances is zero at every observable moment, and no sequence of concurrent requests may produce a state where someone has paid more than they owed. This goal outranks throughput, latency, and feature count, and it is the reason for §4.2 (group-level locking), §4.3 (derived balances), and §4.4 (integer money).

G2. Correctness must be demonstrable, not merely asserted. A design that is correct but unprovable is treated as unfinished. This drives invariant-based testing against real PostgreSQL rather than mocks, the deterministic lock tests in §7.4, and the decision to keep the concurrency model simple enough to reason about exhaustively (§4.2).

G3. Retries must be safe. Clients time out and users press buttons twice; neither may create duplicate financial records, and neither may silently discard data the user believes was saved. This drives §4.5 (idempotency with request hashing).

G4. The database is a participant in correctness, not a passive store. Every application-level rule that can be expressed as a constraint is also expressed as a constraint, so a service-layer bug degrades to a rejected write rather than corrupt accounting (§4.9).

G5. Explicitness over convenience. SQL is written, not generated; transaction boundaries and locks are visible in the code that depends on them. Accepted cost: more mapping code (§4.10).

G6. The system explains its own outputs. A balance can be decomposed into the ledger entries that produced it, with the decomposition verified against the balance at runtime (§11); a plan states its strategy, whether its optimality is proven, and what each transfer costs in new payment relationships (§10). Claims the system cannot prove — greedy optimality, for one — are never made.

Accepted costs. Throughput within a single group is deliberately sacrificed (§4.2). Balance reads are recomputed rather than cached (§4.3). The access model is weak by construction (§4.7). These are consequences of the ranking above, not oversights.

3. System overview

3.1 Components

Component	Responsibility
REST controllers (*Controller in feature packages)	Parse and shape-validate HTTP requests, delegate, map results to JSON
Page controller (GroupPageController)	Same, for server-rendered HTML; holds no business logic
Services (*Service)	Business rules, transaction boundaries, locking, idempotency
Repositories (*Repository)	SQL execution, row mapping, explicit row locking
settlement.plan package	The optimization engine: SettlementPlanner dispatch, GreedyOptimizer, ExactPlanSearch, PlanInvariants; pure functions, no Spring, no I/O (§10)
Error advices (GlobalExceptionHandler, PageExceptionHandler)	Exception → status code + {code, message} or HTML
PostgreSQL	Durable state, constraint enforcement, write serialization

3.2 Data flow: completing a settlement, end to end

The settlement path exercises every mechanism in the system. `SettlementController.complete` receives `POST /api/v1/groups/{groupId}/settlements?token=...` with an `Idempotency-Key` header and a `CompleteSettlementRequest` body, and calls `SettlementService.complete`, which is annotated `@Transactional`. Inside that transaction:

1. `GroupService.requireGroup` loads the group and compares the supplied token to the stored one with `MessageDigest.isEqual`; a mismatch throws `ForbiddenException`.
2. `IdempotencyKeys.validate` rejects a blank or over-long key.
3. `RequestHashes.of` computes a SHA-256 over payer, recipient and amount.
4. `GroupRepository.lockById` executes `SELECT id FROM groups WHERE id = ? FOR UPDATE`. From here until commit, no other expense or settlement in this group can proceed (§4.2).
5. `SettlementRepository.findByGroupIdAndIdempotencyKey` looks for a prior application of this key. If found with the same hash, the stored settlement is returned and the request ends (HTTP 200); if found with a different hash, `ConflictException("IDEMPOTENCY_CONFLICT")` (§4.5).
6. Amount and participant-distinctness are validated.
7. `BalanceRepository.computeBalances` recomputes every balance in the group **inside the lock** — this is what makes the following check sound.
8. The payer must owe at least the amount and the recipient must be owed at least the amount, or `ConflictException("SETTLEMENT_EXCEEDS_DEBT")`.
9. `SettlementRepository.insert` writes the row and returns the database-assigned `created_at`.
10. Commit releases the lock.

Any exception unwinds the whole transaction, so a failure between steps 9 and 10 leaves no row (§8, I7).

```
sequenceDiagram
    participant C as Client
    participant S as SettlementService
    participant G as GroupRepository
    participant B as BalanceRepository
    participant R as SettlementRepository
    C->>S: complete(groupId, token, key, request)
    S->>S: requireGroup via GroupService (token check)
    S->>G: lockById (SELECT ... FOR UPDATE)
    S->>R: findByGroupIdAndIdempotencyKey
    alt key already applied
        R-->>S: existing settlement
        S-->>C: 200 (hash matches) or 409 (hash differs)
    else new request
        S->>B: computeBalances (inside the lock)
        B-->>S: balances
        S->>R: insert
        S-->>C: 201
    end
end
```

Sequence diagram (mermaid notation)

3.3 Read path

Reads take no lock. GET /balances runs the single aggregate of §4.3; GET /activity runs one UNION ALL over expenses and settlements. Both are single statements and therefore internally consistent by construction. The plan and explanation endpoints read several things — balances, relationships, Pix keys, the ledger revision — so each runs as one REPEATABLE_READ transaction: every read sees the same snapshot, and a plan's stamped revision genuinely describes the state it was derived from, even while writers are committing.

4. Architectural decisions

4.1 Module boundary scheme: feature packages over layer packages

Context. The system has six domain concepts (group, participant, expense, balance, settlement, activity) and three layers (controller, service, repository).

Decision. Packages are named for features; layers are expressed by class-name suffix within a package. `common` holds cross-cutting concerns (error contract, shared validators); `web` holds the browser UI.

Alternatives rejected. Layer-first packaging (`controllers/`, `services/`, `repositories/`) was rejected because every change to a feature would touch three distant directories, and because it provides no encapsulation boundary — everything is public to everything. A hexagonal/ports-and-adapters arrangement was rejected as disproportionate: there is exactly one inbound adapter family (HTTP) and one outbound (JDBC), so the indirection would buy nothing.

Consequences. A contributor adding a feature creates one package (§5). The cost is that cross-feature reuse is by direct dependency — `ExpenseService` imports `GroupRepository` and `ParticipantRepository` — so package dependencies form a graph rather than a strict hierarchy. The graph contains one deliberate cycle: `group` and `participant` depend on each other (`GroupService` inserts the creator participant; `ParticipantService` uses the group's token check and lock), reflecting that a group cannot exist without its creator. `common` is a sink — it imports nothing from the application.

4.2 Concurrency: one lock per group, held by every write

Context. Balances are derived from expenses, expense shares, and settlements (§4.3). A settlement is only valid relative to the balances at the moment it commits (G1). Under PostgreSQL's default `READ COMMITTED`, a transaction that validates against a snapshot can be invalidated by a concurrent commit before it inserts.

Decision. Every write path — `GroupService.create` excepted, since the group does not yet exist — acquires `SELECT id FROM groups WHERE id = ? FOR UPDATE` via `GroupRepository.lockById` before reading anything it will validate against: `ExpenseService.create`, `SettlementService.complete`, and `ParticipantService.add`. Reads take no lock.

Alternatives rejected.

- *Locking the participant rows involved.* Rejected because it is unsound: an expense insert changes balances without touching any participant row, so a concurrent expense could commit between a settlement's validation and its insert, breaking invariants I4 and I5 (§8). This is the central concurrency argument of the system.
- *SERIALIZABLE isolation.* Rejected because it converts the failure mode into serialization errors that every caller must be prepared to retry, and because the group lock already makes `READ`

COMMITTED sufficient — all validation happens after the lock is held and inside the same transaction.

- *Optimistic concurrency with a version column on the group.* Rejected as equivalent in throughput to the pessimistic lock for this workload while adding a retry path to test and reason about (G2).

Consequences. All writes within one group serialize. For a five-person dinner this is irrelevant; for a hypothetical thousand-member group it would be a bottleneck. The benefit is that invariants I4, I5 and I8 (§8) have one-paragraph proofs. Note that `ParticipantService.add` takes the lock for explicitness: its inserts were already serialized against other writers as a side effect of the `group_id` foreign key taking `FOR KEY SHARE` on the group row, but that is an invisible property that a schema change could silently remove.

Deadlock analysis. The only conflicting lock mode taken by application code is `FOR UPDATE` on a groups row, and every path takes it before any insert. All participant-row locks are the `FOR KEY SHARE` locks that foreign-key checks acquire, which are mutually compatible. No lock-order cycle is constructible.

4.3 Balances are derived, never stored

Context. Every operation needs current balances; a settlement needs them to be correct.

Decision. No balance column exists. `BalanceRepository.computeBalances` runs one aggregate with four `UNION ALL` legs — expenses paid (+), expense shares assigned (−), settlements sent (+), settlements received (−) — left-joined onto participants so a participant with no activity yields zero via `COALESCE`.

Alternatives rejected. A stored balance updated on each write was rejected because it introduces a second source of truth that can drift from the ledger, and drift in an accounting system is the failure mode with the worst blast radius (G1). A materialized view was rejected because it must be refreshed, which reintroduces staleness inside the settlement transaction, where staleness is precisely what must not happen.

Consequences. Balance reads cost a scan of the group's expenses, shares and settlements — acceptable at the intended scale, and the composite unique constraints provide usable indexes. The zero-sum invariant (I1) becomes a property of the query rather than something application code must maintain. The upper bound on amounts (§4.4) exists partly to keep this aggregate representable.

4.4 Money as integer centavos with an explicit cap

Decision. All monetary values are `long` in Java and `BIGINT` in PostgreSQL, denominated in centavos. Floating-point types never touch a monetary value: `spring.jackson.serialization.accept-float-as-int=false` makes a fractional JSON amount a 400 rather than a silent truncation. `Money.MAX_AMOUNT_CENTS` caps any single amount at 10^{12} centavos, enforced in services and by `CHECK` constraints.

Alternatives rejected. `BigDecimal` was rejected as unnecessary — with a fixed currency and no fractional centavos, integers are exact and cheaper. Unbounded `BIGINT` was rejected after a

demonstrated failure: `SUM(bigint)` returns numeric in PostgreSQL, which does not overflow, so two `Long.MAX_VALUE` expenses produced an aggregate that `ResultSet.getLong` could not read, permanently breaking every balance read for that group.

Consequences. A group cannot record an expense above R\$ 10 billion. The cap must stay consistent between `Money.MAX_AMOUNT_CENTS` and the schema `CHECKS`; they are separate declarations of the same number (§8, A3).

4.5 Idempotency: key plus request hash

Context. Clients retry after timeouts (G3); browsers resubmit forms and offer a back button.

Decision. Expense and settlement creation require an Idempotency-Key, unique per group by constraint. Each row also stores `request_hash`, a SHA-256 over the semantic content of the request (for expenses: description, payer, total, and shares sorted by participant id, so ordering is not significant). On replay: same key and same hash returns the stored record with HTTP 200; same key and different hash returns 409 `IDEMPOTENCY_CONFLICT`.

Alternatives rejected. Key-only idempotency — the earlier design — was rejected once the browser UI existed: editing a form after submission and resubmitting reuses the key with different content, and returning the original record reports success while discarding the user's correction. Silent data loss is worse than a conflict the user can act on. Hashing the raw request bytes was rejected because whitespace or share ordering would produce spurious conflicts.

Consequences. The hash must cover exactly the fields that change meaning; a field added to a request record without being added to the hash creates a silent hole (§8, A4). In the browser, every form POST additionally follows Post/Redirect/Get, so a refresh re-issues a GET rather than a write; the idempotency key handles the retry and back-button cases PRG cannot.

4.6 Settlement plans: derived, strategy-selectable, never stored

Decision. A settlement plan is a pure derivation — current balances, plus a strategy, plus optional constraints — computed per request and never persisted. Three strategies exist (`GREEDY`, `MIN_TRANSFERS`, `RELATIONSHIP_AWARE`); §10 specifies each one. Every plan is stamped with the group's ledger revision (the count of accounting entries), so a caller can tell whether the ledger has moved since the plan was generated.

Alternatives rejected. Storing suggestions as obligations was rejected because any new expense invalidates them, and stale stored obligations would then need invalidation logic — complexity that on-demand generation removes entirely. A single hardcoded algorithm — the original design — was rejected once it was clear that several valid transfer graphs settle the same balances and differ in ways users care about (ADR 0008); baking one choice in hides a decision the group should own.

Consequences. Suggested payments have no identity, so completing one requires the client to submit payer, recipient and amount. A suggestion can go stale between render and submission; the settlement path re-validates under the group lock and rejects with 409 rather than trusting the client. Every strategy is deterministic for a given input, so equal ledger states produce byte-equal plans.

4.7 Access control: invite token, not authentication

Decision. A group is reachable by a 256-bit SecureRandom token, base64url-encoded, compared in constant time. The REST API takes it as a `?token=` query parameter; the browser UI exchanges it once for an `HttpOnly, SameSite=Lax` cookie scoped to `/g/{groupId}` and then uses token-free URLs.

Alternatives rejected. User accounts with passwords or OAuth were rejected as disproportionate for the MVP: they multiply scope (registration, recovery, sessions, per-participant authorization) without exercising the accounting core this system exists to demonstrate. Keeping the token in the URL for the browser was rejected because it lands in history, Referer headers, platform access logs, and chat-app link-preview fetches.

Consequences. Anyone with the link has full read and write access to that group, including other participants' Pix keys. There is no identity, so the system cannot attribute an action to a person or authorize per participant. `SameSite=Lax` is the only CSRF defense; that is sufficient only because the cookie is the sole credential and there is no other session state.

4.8 Append-only accounting

Decision. Expenses and settlements are never updated or deleted. Corrections are made by recording an opposing entry.

Alternatives rejected. Editable expenses were rejected because an edit changes historical balances, which can retroactively invalidate settlements that were valid when they were made — the resulting reconciliation problem is larger than the feature is worth here.

Consequences. No PUT/PATCH/DELETE exists on accounting rows. A mistaken settlement cannot be reversed by an opposite settlement — invariants I4/I5 reject it, because the erroneous payer no longer owes anything — so the correction is a compensating *expense* (paid by the wrongly-recorded recipient, with the full share assigned to the wrongly-recorded payer), which restores both balances. A participant's Pix key also cannot be edited (\$9), which is a real gap rather than a considered decision — it is a consequence of never having built a participant-mutation path.

4.9 Errors: stable machine codes, localized messages, database as backstop

Decision. Services throw `ApiException` subclasses carrying a stable English code (`BadRequestException`, `ForbiddenException`, `NotFoundException`, `ConflictException`). `GlobalExceptionHandler` maps code → HTTP status and resolves the code to a Brazilian Portuguese message via `MessageSource`; `PageExceptionHandler` does the same for HTML. Message resolution has **no default**, so a missing bundle key fails loudly in tests instead of degrading silently in production. Constraint violations that escape service validation surface as 409 `CONSTRAINT_VIOLATION`, and only the first line of the database message is logged, because PostgreSQL puts the offending values — potentially Pix keys — on the `Detail:` line.

Alternatives rejected. HTTP status alone was rejected as too coarse for clients to branch on. Localized text as the contract was rejected because it makes translation a breaking change; tests assert on codes and never on message text.

Consequences. Two advices exist. `PageExceptionHandler` is ordered `HIGHEST_PRECEDENCE` and scoped to the web package, so page controllers get HTML; everything else, including framework-level errors raised before a controller is selected (405, 415, unknown path), falls through to the JSON advice.

4.10 Persistence: JdbcTemplate, no ORM

Decision. Plain SQL through `JdbcTemplate`, with `RowMappers` written by hand.

Alternatives rejected. JPA/Hibernate was rejected because the two properties that matter most here — exactly when a lock is taken and exactly what the balance aggregate is — become implicit under an ORM (G5), and because lazy loading and the persistence context introduce failure modes disproportionate to a five-table schema.

Consequences. Mapping is manual and adding a column touches the record, the mapper, and the insert. The upside is that every statement the system issues is visible in the repository classes.

4.11 Schema application: `schema.sql` at startup

Decision. `spring.sql.init.mode=always` applies `schema.sql`, which uses `CREATE TABLE IF NOT EXISTS`.

Consequences — known defect. This only ever initializes an *empty* database: against a database created by an earlier revision, constraint and column changes are silently skipped and startup reports success. Verified behavior, not speculation. Test suites are unaffected because `Testcontainers` provides a fresh database per run. Migration tooling (Flyway) is the fix and is deferred (§9).

4.12 External dependencies that would hurt to swap

Dependency	Why it is load-bearing
PostgreSQL	FOR UPDATE semantics, deferrable constraints, UNION ALL aggregate, and <code>clock_timestamp()</code> are all used deliberately; a different engine would require re-verifying §4.2 and §4.3
Testcontainers	G2 depends on tests running against the real engine; substituting an in-memory database would invalidate every constraint and locking test
Spring Boot 4.1 / Jackson 3	Jackson 3 lives under <code>tools.jackson.</code> , and <i>parser limits are configured through <code>spring.jackson.factory.constraints.read</code></i> . (§4.13)
Thymeleaf	UI only; swapping it would not touch the accounting core

4.13 Request-size limits at the parser

Decision. `spring.jackson.factory.constraints.read.max-document-length` caps the JSON document, and `CreateExpenseRequest.shares` is bounded by `@Size(max = 500)`.

Context. An unbounded `shares` array was demonstrated to exhaust the heap of a 256 MB instance from a handful of concurrent anonymous requests. Bean validation alone is insufficient: it runs *after*

Jackson has already materialized the list. The parser-level cap also covers chunked requests, which have no Content-Length to inspect.

5. File and directory organization

```
splitpix/
├── src/main/java/com/luiz/splitpix/
│   ├── SplitpixApplication.java    entry point
│   ├── PingController.java         GET /api/v1/ping health endpoint
│   ├── common/                     cross-cutting; depends on nothing in the app
│   │   ├── ApiException.java + Bad/Forbidden/NotFound/Conflict subclasses
│   │   ├── ErrorResponse.java      the {code, message} wire shape
│   │   ├── GlobalExceptionHandler.java exception → status + JSON
│   │   ├── IdempotencyKeys.java     key shape validation
│   │   ├── Money.java              MAX_AMOUNT_CENTS
│   │   ├── RequestHashes.java      SHA-256 over request content
│   │   └── Texts.java              name normalization and rejection rules
│   ├── group/                      group creation, invite token, group view
│   ├── participant/                participants, Pix key types and normalization
│   ├── expense/                   expense creation transaction
│   ├── balance/                   balance aggregate, balance explanation (§11)
│   ├── settlement/                settlement transaction
│   │   └── plan/                  the optimization engine (§10): strategies,
│   │                               search core, constraints, plan invariants,
│   │                               relationship graph, plan service + API
│   ├── activity/                  the ledger read: sequences and revision
│   └── web/                        server-rendered UI; no business logic
├── src/main/resources/
│   ├── schema.sql                 all tables and constraints (§4.11)
│   ├── messages.properties        every user-facing string, pt-BR
│   ├── application.properties     datasource, Jackson limits, SQL init
│   ├── templates/                 Thymeleaf: home, group, planos, extrato, erro
│   └── static/                    splitpix.css, copiar.js
├── src/test/java/com/luiz/splitpix/ API tests, plus settlement/plan/ and web/ unit
tests
├── demo.sh                        walk-through; --tecnico adds the optimizer act
├── scripts/build-design-doc.sh    regenerates the design PDF from this file
└── docs/
    ├── design.md                  this document
    ├── adr/                       one record per load-bearing decision
    └── screenshots/               real captures embedded in the README
```

Placement rule for new code. Code that serves one domain concept goes in that concept's package, named <Concept>Controller|Service|Repository or a record for its request/response/result. Code used by two or more feature packages goes in `common` and must not depend on any feature package. Code that exists only to render HTML goes in `web`. A new domain concept gets a new package rather than an addition to an existing one.

Why `common` is a sink. It holds the error contract and validators that every feature needs. Keeping its dependency set empty keeps the one existing cycle (group ↔ participant, §4.1) from spreading: shared utilities can never become a path between features.

Why web is separate from the feature packages. The UI is a client of the services, exactly as the REST controllers are, and the REST API is the primary interface. Keeping page controllers out of the feature packages makes it structurally obvious that removing the UI would not touch the domain, and lets `ExceptionHandler` be scoped by package (§4.9).

6. Core types and classes

6.1 group

Group — a record mirroring the groups row: id, name, invite token, creation instant. Constructed only by `GroupRepository`'s mapper. Never mutated; a group's name and token are fixed at creation.

GroupService — owns group creation and, critically, `requireGroup(groupId, token)`, the single access check every group-scoped operation calls. It exists as one method so the token comparison has exactly one implementation. `create` is `@Transactional` because it writes two tables (group and creator participant) that must appear together (I7). `getView` returns group plus participants in one read-only transaction at `REPEATABLE_READ` — it runs two statements, and under the default `READ COMMITTED` each would read its own snapshot.

GroupRepository — SQL for groups, plus `lockById`, the concurrency primitive of §4.2. `lockById` throws `NotFoundException` when the row is absent rather than returning silently, because a lock that no-ops is worse than one that fails.

`CreateGroupRequest` / `CreateGroupResult` / `GroupView` / `GroupResponse` — request, service result, internal read aggregate, and wire shape respectively.

6.2 participant

Participant — a record mirroring the row, including the optional Pix key type and value. Created by `ParticipantService.add`; never updated (§4.8).

PixKeyType — enum of `EMAIL`, `PHONE`, `RANDOM`. CPF is deliberately absent: storing a national identifier behind a link-shared access model is a privacy liability with no upside for this system. The enum's absence of a CPF constant is what makes a CPF request a 400 rather than a stored value.

PixKeys — normalization, the paired-nullability rule, and per-type shape validation. Lowercases `EMAIL` and `RANDOM` keys because Pix DICT treats both case-insensitively, so without it two participants could hold the same effective key; re-checks length after normalization because lowercasing can lengthen a string. `validateFormat` requires an email shape for `EMAIL`, E.164 for `PHONE`, and a UUID for `RANDOM` — pragmatic shapes rather than DICT verification, enough to catch a key pasted into the wrong field or a CPF smuggled under `RANDOM` (an EVP key is a UUID, so bare digits never match).

ParticipantService — validates, takes the group lock (§4.2), pre-checks duplicate keys for a friendly error, and maps a `DuplicateKeyException` from the losing side of a race to the same `DUPLICATE_PIX_KEY` code so the contract is timing-independent.

6.3 expense

Expense — the row: id, group, payer, description, total, idempotency key, request hash, creation instant. **ExpenseShare** — a participant/amount pair.

`ExpenseService` — the expense creation transaction (§3.2 has the settlement analogue). Validates that the total is positive and capped, that the payer and every share participant belong to the group, that no participant appears twice, that no share is negative, and that shares sum exactly to the total. Returns shares in a canonical order matching the database's ordering so a replay is byte-identical to the original response.

`CreateExpenseRequest` — carries the bean-validation constraints, including the type-argument constraints (`List<@NotNull @Valid ShareRequest>`) that reject null elements, and the size bound of §4.13. Its compact constructor strips the description so validation and storage agree on length.

6.4 balance

`ParticipantBalance` — participant id, display name, balance in centavos. Produced only by `BalanceRepository`; also the input type of the settlement planners.

`BalanceRepository` — holds `BALANCES_SQL`, the four-leg aggregate of §4.3. This is the single most correctness-critical statement in the system.

`BalanceExplanation` / `BalanceExplanationRepository` — the provenance read of §11: the same four legs as rows, plus the runtime check that they sum to the reported balance.

6.5 settlement and settlement.plan

`Settlement` — the row, including status, which the schema constrains to `COMPLETED`. The MVP has no pending state because there is no asynchronous payment verification to wait for; a status column exists so that adding one later does not require a migration of meaning.

`SettlementService` — the transaction described in §3.2.

The `plan` subpackage is the optimization engine, specified in §10. The pure core — `SettlementPlanner`, `GreedyOptimizer`, `ExactPlanSearch`, `PlanInvariants`, `RelationshipGraph`, `SettlementConstraints` — has no Spring dependency and is unit-tested without a container. `SettlementPlanService` supplies the data (balances, relationships, Pix keys, revision) inside one `REPEATABLE_READ` transaction; `SettlementPlanController` exposes GET (strategy as a parameter), POST (constraints in the body) and the compare endpoint.

6.6 web

`GroupPageController` — one controller for all pages; each action is a POST that redirects (§4.5).

`GroupPage` — a record assembling everything one render needs, so the template performs no lookups.

`InviteCookie` — the token-to-cookie exchange of §4.7. `PageFormats` — a Spring bean (referenced as `@pageFormats` in templates) formatting centavos and masking Pix keys. `MoneyInput` — the strict parser from typed reais to centavos; ambiguous forms (mixed separators, zero-led grouping, three decimals) are rejected, never reinterpreted.

6.7 Trivial types

CreateExpenseResult, CompleteSettlementResult, ExpenseResponse, SettlementResponse, ParticipantResponse, GroupResponse, ErrorResponse, AddParticipantRequest, CompleteSettlementRequest, ActivityItem are records that carry data across a boundary and hold no logic beyond static factory mapping.

7. Interfaces and boundaries

7.1 HTTP → service

Controllers pass primitives and request records; services return domain records. Controllers may assume a service either returns a valid result or throws an `ApiException`; services may assume nothing about the caller and re-validate everything that matters. The boundary protects against a controller — or a future second controller, which is exactly what the UI became — bypassing a rule. This is why `ParticipantService.add` re-checks the token rather than trusting that a controller did.

7.2 Service → repository

Repositories expose intent-named methods (`findByGroupIdAndIdempotencyKey`, `lockById`, `computeBalances`) rather than a generic query interface. Services may assume repositories perform no validation and no business decisions; repositories may assume their arguments are already validated. The boundary protects the ability to read every statement the system can issue by opening one file per table.

7.3 `SettlementPlanner.plan(strategy, balances, relationships, constraints)` → `SettlementPlan`

The strictest contract in the system. Preconditions: balances sum to zero; constraints, when present, name only group members and accompany an exact strategy. Postconditions, enforced by `PlanInvariants.verify` before the plan leaves the planner: applying every transfer drives all balances to exactly zero; every amount is strictly positive and within any cap; no transfer has the same participant on both ends; no payer-recipient pair repeats; no forbidden pair appears; total sent equals total owed. Plans are deterministic for a given input. `exact = true` additionally means the plan is provably optimal for the strategy's declared objective (§10.4). Failure modes are typed: oversized group → `UNSUPPORTED_OPTIMIZATION_SIZE`, unsatisfiable constraints → `NO_FEASIBLE_SETTLEMENT_PLAN`, constraints with GREEDY → `INVALID_SETTLEMENT_CONSTRAINT`, and an optimizer bug caught by the validator → `IllegalStateException`, surfacing as a 500 rather than a wrong plan.

7.4 Test seams

`ApiTestSupport` is the shared base for API tests; sharing one Spring context and one container across subclasses is deliberate, since each distinct annotation combination forks a new context and a new PostgreSQL container. `@MockitoSpyBean` on a repository is the seam used to force failures that cannot be produced through the API — rollback tests let the real insert execute and then throw, because a test that mocks the insert away would pass without any transaction at all.

`GroupLockTest` deserves mention as a boundary of a different kind: it holds the group lock inside a paused transaction and asserts the second request has not completed, because a barrier-synchronized race passes in a warm test context even with the lock removed.

8. Invariants and assumptions

Violations here are silent and expensive. Nothing in the type system enforces them.

I1 — A group's balances sum to zero. Guaranteed by the aggregate's structure (§4.3) as long as every leg's sign is correct and every referenced participant belongs to the group. *Breaks if:* a leg's sign is flipped, or a row references a participant of another group. The composite foreign keys make the second case impossible at the database level.

I2 — An expense's shares sum to its total. Enforced in `ExpenseService.validate` only. The database cannot express it. *Breaks if:* the sum check is removed or bypassed by direct SQL; balances then stop summing to zero.

I3 — Every accounting row references participants of its own group. Enforced by composite foreign keys on `(group_id, participant_id)`, which require the `UNIQUE (group_id, id)` constraints on participants and expenses to exist. *Breaks if:* those unique constraints are dropped — the composite keys become undeclarable and the schema silently loses the guarantee.

I4/I5 — No over-settlement in either direction. Enforced in `SettlementService` under the group lock. *Breaks if:* the lock is removed (two individually valid settlements can then jointly over-settle), or either side of the balance check is dropped.

I6 — One idempotency key produces at most one durable record. Enforced by the group lock plus `UNIQUE (group_id, idempotency_key)`.

I7 — A failed transaction leaves no partial state. Depends on `@Transactional` being present on every multi-statement write and on Spring's proxying applying — a self-invocation would silently bypass it. No such call exists today; `requireGroup` is deliberately unannotated so it joins the caller's transaction.

I8 — Concurrent settlements cannot over-settle. A corollary of I4/I5 plus §4.2, and the reason the group lock exists.

I9 — Every emitted settlement plan settles the group. Enforced by `PlanInvariants.verify` on every generated plan, in production, not only in tests: amounts positive, no self-payment, pairs unique, constraints honored, and applying the plan leaves every balance at exactly zero. *Breaks if:* the validator call is removed — an optimizer bug then ships as a payment suggestion instead of a 500.

I10 — A balance explanation sums to the balance it explains. Enforced at runtime in `BalanceService.explain` (§11) under `REPEATABLE_READ`. *Breaks if:* the balance aggregate and the explanation query stop describing the same legs — which is precisely the drift the runtime check exists to catch.

A1 — `clock_timestamp()`, not `now()`, orders the ledger. `now()` is transaction-start time, but writes `serialize` at lock-acquisition time, so a transaction that waited would receive a timestamp earlier than writes it depends on, and the activity feed could show a settlement before the expense that created the debt.

A2 — The database is initialized empty. See §4.11. A database created by an older revision silently lacks newer constraints.

A3 — `Money.MAX_AMOUNT_CENTS` and the schema CHECKs state the same number. Two independent declarations; a change to one without the other produces either a 500 where a 400 belongs, or an uncapped path.

A4 — `request_hash` covers every field that changes a request's meaning. A field added to a request record but not to the hash makes two materially different requests interchangeable under one key.

A5 — Pix key values are never logged. Enforced by convention plus the first-line truncation in `GlobalExceptionHandler`; PostgreSQL puts values on the `Detail:` line of constraint messages.

A6 — Every error code emitted has a bundle key. Message resolution has no default (§4.9), so a missing key is a runtime failure, not a degraded message.

A7 — The exact-search thresholds match measured behavior. Ten nonzero balances uncapped, eight capped, backstopped by a deterministic node budget (§10.6). The numbers came from measurement on adversarial instances; raising them without re-measuring turns a 400 into a slow request.

9. Known gaps and future work

Gap	Status and reason
Schema migrations	CREATE TABLE IF NOT EXISTS cannot evolve an existing database (§4.11). Flyway is the intended fix; deferred because every environment so far is disposable.
Pix key editing	Not implemented. A key typed incorrectly can never be corrected: a wrong key appears in full on every payment instruction, and a participant with no key appears as a recipient with the key shown as absent. This is the most user-visible gap. Participants are not accounting rows, so append-only (§4.8) does not forbid a mutation path — it was never built.
Participant removal	Not implemented; the deferred foreign key on expense_shares deliberately makes a raw delete fail rather than silently corrupt balances.
Group and participant creation idempotency	The browser forms for creating a group and adding a participant carry no idempotency key, and the corresponding endpoints accept none — a back-button resubmission duplicates the participant (the expense and settlement paths are protected).
Historical balance snapshots	"Balances as of ledger entry N" is not implemented natively. The ledger endpoint's dense sequences make external reconstruction possible; a native version would need a cross-table cutoff that is easy to get subtly wrong, and a clean ledger view was judged better than a fragile snapshot feature (§11).
Rate limiting and abuse controls	Absent. POST /api/v1/groups is unauthenticated by design, so a public deployment needs per-IP limits and caps before exposure. The exact-plan endpoints add CPU-bound work per request (bounded by the node budget), which strengthens the case for per-IP limits.
Public hosting	Not deployed. Preconditions: migrations, rate limiting, a declared data-retention posture (Pix keys are personal data), and confirmation that no invite token reaches platform access logs.
Group-existence oracle	GET returns 404 before checking the token, so an unauthenticated caller can distinguish existing from non-existent group ids. Accepted: ids are random UUIDs and existence alone reveals nothing.
Unknown-path responses on page routes	A URL matching no controller returns the JSON error contract rather than an HTML page, because advice scoped to a package cannot apply when no controller was selected (§4.9).
Observability	Only errors are logged. No request id, no metrics, no per-operation success logging.
PARTICIPANT_NOT_FOUND	No such code exists; an unknown participant id is reported as 400 PARTICIPANT_NOT_IN_GROUP. Recorded because the v2.1 document promised a 404.

10. The settlement optimization engine

10.1 The problem, precisely

Given nonzero balances $b_1 \dots b_n$ summing to zero, a settlement plan is a list of transfers (payer, recipient, amount > 0), payer a debtor and recipient a creditor, such that applying every transfer leaves all balances at zero. Many plans exist for the same vector. Two measures distinguish them here:

- **transfer count** — how many payments people must actually make;
- **novel relationship edges** — transfers between participants with no prior financial relationship in the group (§10.5).

Minimizing transfer count is NP-hard: the minimum equals n minus the largest number of disjoint zero-sum subsets the vector can be partitioned into, and recognizing zero-sum subsets embeds subset-sum. So there is a real algorithmic trade: a fast heuristic with no guarantee, or an exact search with a size limit. SplitPix ships both and labels which is which.

10.2 A worked example

Balances, in reais: Ana +500, Bruno +400, Clara −400, Diego −300, Elisa −200 (from two expenses: Ana paid a 500 hotel split by Diego and Elisa; Bruno paid a 400 car rental for Clara).

Greedy pairs the largest debtor with the largest creditor, repeatedly:

Clara → Ana	400	(Ana now +100)
Diego → Bruno	300	(Bruno now +100)
Elisa → Bruno	100	
Elisa → Ana	100	– 4 transfers

Pairing Clara with Ana destroyed the {Bruno +400, Clara −400} component that the expense structure had already created. The true minimum:

Clara → Bruno	400	(one zero-sum component)
Diego → Ana	300	
Elisa → Ana	200	– 3 transfers

Three transfers, and every one follows an existing expense relationship — here MIN_TRANSFERS and RELATIONSHIP_AWARE agree. They do not in general: with balances $A - 600, B - 400, C + 600, D + 400$ and relationships $\{A-D, B-C\}$ only, the two-transfer minimum ($A \rightarrow C, B \rightarrow D$) creates two novel pairs, while the relationship-aware optimum spends a third transfer to get down to one novel pair ($A \rightarrow D$ 400, $A \rightarrow C$ 200, $B \rightarrow C$ 400) — and one is provably the floor, because the related edges alone carry at most 800 of the 1000 owed. Both cases are pinned in SettlementPlannerTest.

10.3 The plan space: basic plans

Every strategy searches (or emits into) the same space: plans in which each transfer either zeroes one participant or saturates a capped edge. These are the basic solutions of the underlying (capacitated) transportation problem. The space is exhaustive in the sense that matters: any feasible plan can be reduced — by canceling flow around cycles until an edge empties or hits its cap — to a basic plan on a subset of its own edges, which therefore has no more transfers and no more novel edges than the original. Searching basic plans only is a completeness argument, not an approximation.

10.4 The strategies

GREEDY (`GreedyOptimizer`). Two priority queues, largest debtor against largest creditor, transfer the smaller magnitude, deterministic tie-break by participant id. $O(n \log n)$, at most $n-1$ transfers, any group size. Reports `exact = false` unconditionally — §10.2 shows why the claim would be false.

MIN_TRANSFERS (`ExactPlanSearch`, transfers-only objective). Depth-first search over basic plans: at each state, for every active (debtor, creditor) pair, transfer $\min(\text{debt}, \text{credit}, \text{cap})$ and recurse. Memoized on the remaining balance vector — the cost to finish from a state does not depend on the path into it. Returns a provably minimal plan (`exact = true`).

RELATIONSHIP_AWARE (same core, lexicographic objective). Cost is the pair (novel edges, transfers) compared lexicographically; **MIN_TRANSFERS** is the degenerate case that charges zero for novelty. The result is exact for the declared objective: fewest transfers between unrelated people first, fewest transfers second.

Determinism, all strategies: pairs are visited in participant order and ties keep the first plan found, so equal inputs give equal plans — which the API relies on for replay-stability and the tests rely on for pinning.

10.5 The relationship graph

Defined in ADR 0010, derived in one query: A–B are related when one paid an expense the other held a positive share of, or a completed settlement exists between them, either direction. Undirected, unweighted; zero shares create no edge; co-sharers of the same expense are not related to each other (the money flowed through the payer). Settling a suggested payment creates a real edge, so a novel pair used once stops being novel.

10.6 Complexity, thresholds, and the node budget

The search is exponential in the worst case. Measured on 200 randomized adversarial instances per configuration (values chosen to make zero-sum subsets rare, all pairs allowed, empty relationship graph — the worst shape), on an Apple-silicon laptop:

Configuration	Worst case
10 nonzero balances, no cap	59 ms
8 nonzero balances, per-transfer cap	136 ms
12 nonzero balances, no cap	exceeds the node budget

Hence the limits: **10 nonzero balances** for exact strategies, **8 with a cap** (a saturating move zeroes nobody, so used edges join the memo key and states repeat far less). Past the limit: 400 `UNSUPPORTED_OPTIMIZATION_SIZE`, never a silent fallback (ADR 0009). A deterministic budget of five million search states backstops the thresholds; identical input explores identical states, so the budget can never fail intermittently for a given request. `PlanSearchPropertyTest.worstCasesAtTheThreshold` keeps boundary instances inside the budget.

10.7 Constraints

Two types (ADR 0011): directed forbidden pairs and a per-transfer cap. Both prune moves inside the search, so satisfaction is by construction; both are re-checked by `PlanInvariants` on the way out. A plan carries at most one instruction per payer-recipient pair, so a cap below a two-person debt is genuinely infeasible — answered with 409 `NO_FEASIBLE_SETTLEMENT_PLAN`, the same verdict the search returns whenever constraints admit no plan. GREEDY plus constraints is a 400: greedy cannot honor them without sometimes reporting false infeasibility, and a constraint honored "usually" is worse than none.

10.8 How the engine is validated

Three independent oracles, none sharing code with the solver:

1. **Partition bound.** For random vectors up to $n = 8$, an independent subset DP computes $n - \max$ zero-sum partition — the closed-form minimum — and the solver's transfer count must equal it (300 seeds).
2. **Exhaustive enumeration.** For random small instances with random relationships, forbidden pairs and caps, a memo-less enumeration of the entire move space computes the true lexicographic optimum — and the true infeasibility verdict — which the solver must match (250 seeds).
3. **Re-application.** Every plan from every strategy is re-applied to the balances in test code (not via `PlanInvariants`, so a checker bug cannot mask a solver bug) and must zero the group.

Plus the pinned counterexamples of §10.2, determinism checks, and threshold boundary tests. In production, every plan passes `PlanInvariants.verify` before it leaves the service (I9).

11. Balance provenance

GET /participants/{id}/balance-explanation decomposes a balance into the ledger entries behind it: the same four legs the balance aggregate sums — expenses paid (+), shares held (–), settlements sent (+), received (–) — as rows in ledger order, zero shares omitted.

The design constraint (ADR 0012) is that the explanation must be the same truth as the balance, not a parallel calculation that can drift. Two mechanisms enforce it. The queries share one REPEATABLE_READ snapshot, and the service sums the entries and compares against the balance aggregate's answer: a mismatch is an `IllegalStateException` — a 500 — because a statement that almost reconciles is worse than no statement. The equality is also asserted over randomized activity in `BalanceExplanationApiTest`, and the browser statement page (`/g/{id}/extrato/{participantId}`) renders the rows with the reconciled total as its closing line.

The ledger view itself (GET /activity) carries a dense 1-based sequence per entry, ordered by (`created_at`, `id`) — trustworthy as serialization order because writes hold the group lock and timestamp with `clock_timestamp()` (A1). The highest sequence is the group's ledger revision, the same number plans are stamped with. Point-in-time balance reconstruction ("balances as of entry N") is not implemented; the ledger endpoint provides the entries to do it externally, and a native implementation is listed in §9.

12. Against a conventional expense-sharing design

The comparison is with the conventional way these systems are built — a mutable balances table and one hardcoded settle-up routine — not with any specific product; competitors' internals are not verifiable from here.

Dimension	Conventional design	SplitPix
Balance storage	mutable balance column, updated per write	derived from the ledger on every read; zero-sum is a query property (§4.3)
Corrections	edit or delete the expense	append a compensating entry; history is never rewritten (§4.8)
Concurrent writes	last-write-wins or row-level accidents	serialized per group by an explicit lock, with invariants I4/I5/I8 checked inside it (§4.2)
Retries	duplicate rows or client-side dedup	idempotency key + request hash; replay is byte-identical, mutation is a 409 (§4.5)
Settlement output	one algorithm, usually greedy, often labeled "simplify debts"	three strategies with declared objectives; exactness is a tested claim, not a label (§10)
Constraints	none	forbidden pairs and per-transfer caps, enforced exactly or refused (§10.7)
Explainability	a number	any balance decomposes into ledger entries that provably sum to it (§11); any plan states its strategy, guarantee and novel-edge cost
Payments	some execute payments	payment instructions only, never execution or verification (§1.4)
Testing	mocked persistence	real PostgreSQL, deterministic lock tests, seeded property tests, independent brute-force oracles (§7.4, §10.8)

The right-hand column is the project's identity in one table: a correctness- first ledger under an ordinary expense-splitting UI, with the settlement graph treated as a decision rather than a side effect.

Appendix A: glossary

Balance — a participant's net position in centavos: positive means the group owes them, negative means they owe the group. Always derived (§4.3).

Basic plan — a settlement plan in which every transfer zeroes a participant or saturates a capped edge; the space the exact search enumerates (§10.3).

Centavo — 1/100 of a Brazilian real; the atomic unit of all monetary values here.

DICT — the Brazilian Central Bank's directory that maps a Pix key to an account. SplitPix does not query it; it is the reason email keys are treated case-insensitively.

Group lock — the `SELECT ... FOR UPDATE` on a groups row that serializes all writes within one group (§4.2).

Idempotency key — a client-supplied identifier, unique per group, that makes a retried write produce at most one record (§4.5).

Invite token — the 256-bit secret that grants access to a group; the system's only credential (§4.7).

Ledger revision — the count of a group's accounting entries (expenses plus completed settlements). Append-only, so it only grows; plans are stamped with the revision they were derived from.

Novel relationship edge — a transfer between two participants with no prior financial relationship in the group (§10.5); the quantity `RELATIONSHIP_AWARE` minimizes first.

Pix — Brazil's instant payment system. SplitPix coordinates around it but performs no transfers.

Pix key — an email address, phone number, or random UUID that identifies a payee in Pix.

Relationship graph — the undirected graph over participants derived from the ledger: an edge means a shared expense (payer to positive-share holder) or a completed settlement (ADR 0010).

Settlement — a recorded assertion that one participant paid another; append-only and always COMPLETED.

Settlement plan — a derived list of payment instructions that would zero every balance; carries its strategy, exactness claim and metrics (§10). Never stored.

Strategy — one of `GREEDY`, `MIN_TRANSFERS`, `RELATIONSHIP_AWARE`: the declared objective a plan was computed under (§10.4).

Share — the portion of an expense assigned to one participant; shares of an expense sum to its total (I2).

Suggested payment — one transfer of a settlement plan; generated on demand, never stored (§4.6).